

Problem Statement & Product Discovery Document

Aaryajanani — Social Media Content Operations

Based on the real-world workflow, incidents, pain points, and spoken problem-discovery discussion.

Status: Problem discovery completed for the current scope

Important boundary: This document intentionally does NOT prescribe the final product solution, UI, technology stack, or feature set. It documents the problem as observed.

1. Executive Summary

Aaryajanani is an NGO with an active social-media presence, including Instagram and YouTube. Its current social-media operation is handled by a small team of four people whose combined workload is closer to the work normally distributed across a larger team.

The core problem is not simply that WhatsApp, Google Sheets, or Google Drive are imperfect tools. The deeper problem is that the lifecycle and memory of a content project are fragmented across multiple tools, conversations, files, and individual memory. There is no single persistent, connected record that reliably represents what a project is, how it evolved, what was decided, what was approved, what was published, and how it performed.

As a result, people repeatedly perform manual coordination and reconstruction work: asking for task status, finding files, determining which revision was approved, remembering past experiments, locating final assets, and reconnecting a published post with the project that produced it.

2. Discovery Scope and Source

This document consolidates the user's real-life observations of Aaryajanani's current social-media workflow and the subsequent one-question-at-a-time problem-grilling discussion. The examples below are treated as observed scenarios, not hypothetical product requirements.

Primary source context:

- Real organization: Aaryajanani (NGO).
- Current social-media presence: Instagram and YouTube.
- Current team size: 4 people.
- Observed workload: approximately the workload of a much larger team.
- Current working tools mentioned: Google Meet, Google Sheets, WhatsApp, Google Drive, and social-media platforms.

3. Actors and Responsibilities

Actor	Primary responsibilities	Observed workload	Key problem exposure
-------	--------------------------	-------------------	----------------------

A — Team Lead / Lead	Maintains the team; hires members; scheduling; publishing; HR; task assignment; coordination; growth planning and experimentation; alignment with organizational principles.	Very broad. Carries management, HR, operational, scheduling, publishing, and strategic responsibilities.	Cannot reliably track what C/D completed or are currently doing; has to call them for status because the Sheet is not reliably maintained.
B — Content / Social Media	Research; information gathering; script writing; appearing in/shooting videos; providing raw footage, voiceovers and inputs; planning reels/videos; also publishing content.	High content-production and coordination load.	Must chase input delivery; participates in review/approval; can struggle to identify the correct approved version during publishing.
C — Editor / Designer	Primarily video editing; some design; scheduling; performance analysis; future-post planning; experiments; publishing.	Carries substantially more video-editing and operational work than D.	Extra manual status updates; fragmented files/reviews; must repeatedly communicate draft and revision state.
D — Editor / Designer	Posters, carousels, data gathering for those assets, and some video editing.	More focused on static design work; posters/carousels are relatively infrequent.	Same general tracking and file-organization problems apply when work moves through the shared workflow.

4. Content Context

The team's social-media content includes, among other things, ashram testimonials and stories related to pre-pregnancy and post-pregnancy themes drawn from Puranic/traditional material.

5. Current End-to-End Workflow

The observed workflow for a new video concept is approximately:

1. The team meets, often through Google Meet, to brainstorm a new video concept.
2. The discussion is collaborative; the concept is not treated as a one-person approval decision. The team discusses it and agrees to work on it.
3. The concept is noted in a Google Sheet.
4. Required inputs for the editor/designer are discussed with B, often orally.
5. C or D must separately ask B when the required inputs will be delivered. There is currently no reliable tracking of that dependency.
6. Once inputs are received, C or D works on the content.
7. The first draft is posted in the group, historically via WhatsApp, and the Sheet is manually updated to indicate 'First Draft Delivered.'
8. A and B review the draft and comment on it.
9. A, B, and the responsible editor discuss the feedback. The editor can challenge feedback or explain why a review point may be incorrect. The process is collaborative rather than purely top-down.
10. The team reaches a conclusion on the changes.
11. The editor works on revisions and manually updates the Sheet to indicate that changes are in progress.
12. The revised work is submitted again for review and discussion.

13. Eventually A or B gives final approval; A generally has higher authority when a final decision is needed.
14. The project is manually marked as completed.
15. The completed project is stored in Google Drive.
16. On the scheduled date, the scheduler/publisher retrieves the completed file from Drive and publishes it.
17. If the final file cannot be found, the scheduler asks the editors or group where the project is and who worked on it.
18. After publishing, the published social-media post and its performance are not currently connected to a durable project record.

6. Where Information Currently Lives

Information / event	Current location	Observed consequence
Concept / brainstorming	Google Meet and/or later noted in Google Sheet	Decision context can be separated from the eventual project history.
Task / status	Google Sheet	Requires manual maintenance; often becomes stale or unclear.
Input requests and delivery timing	Oral conversation / WhatsApp	No reliable tracking of when inputs are expected or received.
Drafts	WhatsApp/group	Versions and context can become buried.
Review comments	WhatsApp/group	Feedback is not inherently tied to a persistent project/version state.
Revision discussion	WhatsApp/group	People must reconstruct which version a comment refers to.
Approved version	Implicitly determined from conversation	A reviewer may know a version was approved, but the approval/version relationship is not clearly recorded.
Raw/final assets	Google Drive	Drive structure is not consistently organized or relational.
Schedule	Team process / scheduler's knowledge	Publisher may need to search for the correct final asset.
Published post	Instagram/YouTube	Not currently linked back to the project.
Performance	Social platforms / separate analysis	Not currently connected to the originating concept/project in one persistent record.

7. Core Problem Statement

Aaryajanani's content team does not have a single, persistent, connected record of a content project from concept through production, review, completion, publication, and performance.

Instead, project information is fragmented across meetings, spreadsheets, chat conversations, file storage, and individual memory. The team therefore has to perform repeated manual coordination and reconstruction work to answer basic operational questions such as: What is this project? Who is responsible? What version are we on? What feedback was given? What was approved? Where is the final file? When is it being published? What was actually published? And how did it perform?

The resulting inefficiency is not merely a lack of convenience. It creates uncertainty at critical transition points in the content lifecycle and makes the organization's project history difficult to preserve, retrieve, and understand.

8. Pain Points

Pain point	Observed manifestation
Fragmented project context	A single project is spread across Google Meet, Sheets, WhatsApp, Drive, and social platforms.
Manual status maintenance	Editors must update the Google Sheet after doing the actual work; this is viewed as extra work.
Unclear status visibility	The lead cannot reliably know what C/D have completed, started, or are currently doing without asking them.
Poor shared readability	Even when C updates the Sheet, other team members may not understand what C has written.
Untracked dependencies	C/D have to ask B when inputs will be delivered; there is no dependable record of the dependency.
Version confusion	Multiple revisions can exist, but the relationship between a version and its approval is not reliably recorded.
Approval ambiguity at the artifact level	A/B can approve a particular version, but the approval is buried in conversation rather than represented as durable project state.
Review context loss	Comments, counter-comments, and negotiated decisions are mixed into group conversation.
File discovery failures	The scheduler may not find the completed project in Drive and must ask the team.
Ownership discovery failures	There has been an instance where the scheduler forgot who worked on a project and had to ask the group.
Unorganized asset relationships	Inputs, drafts, revisions, and finals are not reliably connected in Drive.
Experiment memory loss	A previously discussed experiment was later forgotten by A/B, forcing C to explain the proposal again.
Weak project history	There is no durable record connecting concept → work → review → final → published post → performance.
Post-publication disconnect	The social-media post link is not currently tied back to the project/concept.
Performance disconnect	Performance of a published video is not persistently tied to the originating concept/project.

9. Serious Real-World Scenarios

Scenario 1 — Wrong/uncertain version at publishing time

Project 1 went through multiple revisions. C used a naming convention such as projectName_versionNo. C sent Version 3; B reviewed it and indicated that it looked good. Later C sent another version. A commented elsewhere. The project discussion and files were distributed across WhatsApp. No durable record clearly stated 'Version 3 = B approved; A pending' or which later version superseded it. When the scheduled publishing time arrived, B could not confidently determine which version was approved. The relevant file and comments were also difficult to recover from WhatsApp.

Operational consequence: publishing became uncertain at the exact moment the content needed to be released.

Scenario 2 — Lead must manually interrogate the team for status

A needed to know what C and D had completed, what they were currently working on, and what tasks had been initiated. The Google Sheet could not reliably answer these questions because team members were not consistently maintaining it. A therefore had to call C/D and ask them directly.

Operational consequence: the lead spends time reconstructing team status instead of obtaining it directly from the workflow.

Scenario 3 — Experiment proposal is forgotten

The entire team discussed creating a new type of video as an experiment. Later, C asked A and B about the experiment. Neither remembered the proposal. C had to explain the experiment again.

Operational consequence: organizational knowledge created collaboratively did not remain reliably available after the meeting.

Scenario 4 — Scheduler cannot locate final asset or owner

On the scheduled publishing date, the scheduler needs the completed video from Google Drive. When the file cannot be found, the scheduler asks the editors. There has also been a case where the scheduler forgot who had worked on a project and had to ask the group who created it and where the file was.

Operational consequence: publishing depends on human memory and ad-hoc group queries.

10. Root Causes and Contributing Factors

The following are problem hypotheses grounded in the observed workflow. They are intentionally separated from product features.

- Multiple systems contain different pieces of the same project's state.
- The system of record for a project is not clearly defined.
- Manual data-entry steps are disconnected from the actual work, so maintaining status feels like administrative overhead.
- Important decisions are made in conversation but are not consistently converted into durable project state.
- File storage is organized as storage, not as a connected representation of project relationships.
- Dependencies such as 'B must deliver inputs to C' are communicated but not persistently tracked.
- Review is collaborative and iterative, but the final relationship between comments, revisions, and approved versions is not explicitly preserved.
- Publishing is operationally dependent on finding the right final artifact, but the final artifact is not reliably connected to the project and its schedule.
- Post-publication information is disconnected from the originating project, making the content lifecycle incomplete.
- The team is small relative to its workload, increasing the cost of every manual coordination step.

11. What the Problem Is NOT

- It is not simply 'WhatsApp is bad.' WhatsApp is useful for fast discussion and collaborative review.
- It is not simply 'Google Sheets is bad.' The Sheet becomes problematic because it requires separate manual maintenance and does not reliably represent the full project state.
- It is not simply 'Google Drive is bad.' Even a perfectly organized Drive would not solve fragmented discussions, approvals, scheduling, and performance history.
- It is not primarily an approval-hierarchy problem. A/B/C can openly discuss feedback; A generally has higher final authority, but review is collaborative.

- It is not merely a lack-of-people problem. The team is indeed overloaded, but the observed failures are specifically related to coordination, traceability, and information fragmentation.

12. Required Organizational Memory (Problem Definition, Not Solution)

From the discovery discussion, the team considers the following information especially important to preserve for a specific content project:

- The concept that was discussed and agreed upon.
- The fact that the concept became a real video/project.
- The assigned work and project lifecycle.
- Changes and revisions made during production.
- Comments, feedback, and review history.
- The final/published video or social-media post link.
- Performance of the published concept/video.

The identity of the individual responsible for the work is useful operational information, but in the discovery discussion it was explicitly judged not to be the most critical long-term piece of information. The project/concept, published content, and performance were considered more important to preserve. This does not mean ownership is unimportant during active work.

13. Project Lifecycle as Observed

The emerging mental model is similar to a persistent GitHub issue: a project/concept should remain identifiable as one continuous unit while its state changes.

Observed lifecycle:

19. New concept / proposal
20. Collaborative discussion
21. Agreement to pursue the concept
22. Assignment / responsibility
23. Input collection
24. Production
25. First draft
26. Review and discussion
27. Revision
28. Final review
29. Completion
30. Scheduling
31. Publication
32. Performance measurement

The same persistent project record is expected conceptually to support different kinds of work/events, including new concepts, new videos, revisions, experiments, and project completion, rather than forcing every activity into an unrelated standalone record.

14. Important Relationships Discovered

Relationship	Why it matters
Concept ↔ Project/video	The published work should remain traceable to the original concept.
Project ↔ assigned work	During active production, the team needs to know who is handling the work.
Project ↔ drafts/revisions	The evolution of the work needs to remain understandable.
Project ↔ comments/reviews	Review reasoning and negotiated decisions need context.
Project ↔ final approved artifact	Publishing depends on knowing which artifact is final.
Project ↔ scheduled publication	The publisher needs to know what is due and what asset should be used.
Project ↔ published social-media post	The team should be able to trace a project to the actual published result.
Project/concept ↔ performance	Performance needs to be understood in relation to what the team created.
Experiment ↔ resulting content	An experiment should be traceable to the content produced and the eventual learning.

15. Operational Consequences

- Time is lost to asking people for information that should be readily available.
- Publishing can become risky because the correct final version is not immediately identifiable.
- Managers cannot reliably see team workload and current work without manual follow-up.
- Review cycles become harder to reconstruct when multiple revisions occur.
- Past decisions and experiments can be forgotten.
- New or less-informed team members have difficulty understanding historical context.
- Files become detached from the project context that explains what they are.
- The organization cannot easily connect content production to published outcomes and performance.
- Operational knowledge remains partly dependent on individual memory.

16. Open Questions and Unverified Areas

The problem discovery is strong, but the following should remain explicitly unverified until further observation/interviews:

- Exact frequency of each failure scenario.
- Exact number of content projects per month.
- Exact average number of revisions per project.
- Whether every project follows the same workflow or whether some content types differ.
- Exact location and method currently used for scheduling information.
- How performance metrics are currently collected and analyzed.
- Whether there are additional pain points for the NGO beyond social-media operations.
- How permissions should vary by role in practice. The discussion indicates that access can be controlled by an admin/team lead, but detailed permission rules were not fully specified.

17. Problem Discovery Conclusion

The strongest evidence points to a coordination and project-memory problem caused by fragmented information, not to a single bad tool.

Aaryajanani's team has a collaborative content-production process, but the process crosses several disconnected systems. The team's actual work, discussions, decisions, files, status, scheduling, publication, and performance are not represented as one continuous project history. This forces people to manually bridge gaps between systems and repeatedly reconstruct state.

The central problem to preserve for future product work is therefore:

“The team can create and discuss content collaboratively, but the organization cannot reliably preserve and retrieve the complete connected history and current state of that content from concept to performance.”

18. Discovery Artifact: One-Line Problem Map

Fragmented information → manual coordination → uncertain project state → file/version confusion → publishing friction → lost decisions/experiments → disconnected publication/performance → repeated reconstruction work.

19. Explicit Boundary Before Solution Design

No final solution has been selected in this document. The GitHub-style issue/repository analogy is recorded as the user's mental model for persistent project history, not as a committed implementation decision. The next product-design phase should test whether that model actually fits all observed workflows before defining screens, features, database schemas, or architecture.